

Token based Auth

- After login, the server gives the client a token. client sends that token with every request to prove identity. Instead of remembering the user on the server (sessions), client carries proof of authentication.
- * Traditional session-based authentication stores session data on the server.
- But when system scales horizontally, session storage becomes difficult.

Solution

Token based auth solves this by:

- Removing server side session storage.
- Allowing stateless APIs, making scaling easier

Working

- Login - user sends email+pass → server verifies credentials and generates a token & sends it back to client.
- Client stores token in memory, HttpOnly cookies.
- Client sends token in header & server verifies & grants access (No DB lookup if JWT).

Security Consideration

- 1) Always use HTTPS (token must never travel unencrypted)
 - 2) Store tokens safely (Use HttpOnly secure cookie)
 - 3) Short expiry
 - 4) Token revocation strategy.
- JWT is stateless. To revoke:

- Maintain blacklist
- Rotate signing keys

Note:- JWT is more scalable than session not automatically secure.

Access Control Lists & Rule Engines

- ACL is a list that defines who can access what & what actions they can perform.

- Suppose we have 100 users & 500 documents.
Each document may have: owner, editor, viewer.

ACL would store rules like:

- User 1 $\rightarrow$ read $\rightarrow$ doc 45
- User 2 $\rightarrow$ edit $\rightarrow$ doc 45

* ACL is used in:-

- File systems
- Cloud storage (S3 bucket policies)
- Database row-level security
- Document management system.

Rule Engines: A system that evaluates logical conditions to decide whether something is allowed.

- Instead of storing static permissions, it evaluates rules dynamically.

* These are needed because modern apps need:

- Dynamic authorization
- Policy-driven system.
- Context based access

* Rule engines allow flexible policies, changing rules without changing code, scalable permission logic

Real world System

Large system combine:-

- Authentication + ACL + Rule Engine

EX:- Banking system

ACL $\rightarrow$ user can access account

Rule engine $\rightarrow$ user can transfer only if $bal > x$.

Rate Limiting

- Restricting how many requests a user (or IP) can make in a certain time.

Why :-

- Attackers can brute force password.
- user can overload your API.

It can be applied:

- Per IP address
- Per API Key
- Per user ID
- Per route

Algorithms

- i) Fixed window :- 100 requests / min
- ii) Sliding window :- counts requests in a moving time window
- iii) Token bucket (common) :- System gives tokens at fixed rate.
 - Each request consumes one token.
 - If no tokens left $\rightarrow$ blocked.

DDoS Protection

- Distributed Denial of Service.
- Millions of machine flood your server with traffic at the same time to make your system unavailable.
- Rate limiting works at application level but DDoS attacks happen at network level.

Working

- 1) CDN (content delivery network) absorb traffic before it reaches your server.
- 2) Web App Firewall filters bot traffic, malicious pattern.
- 3) Load balancers distribute traffic across multiple servers.
- 4) Auto scaling (servers are added automatically)
- 5) Network level filtering (Block traffic at edge level).

Distributed Rate Limiting

- * Enforcing rate limits consistently across multiple servers. All servers must share same request count.

* Need

Modern system scale horizontally, runs in containers & uses microservices.

Without DRL

- Limits can be bypassed, stability of system reduces.
- All servers must use shared central storage to track request counts.
- Instead of storing counters in memory, we store them in:
 - Redis
 - Database (less common)
 - Distributed cache

Why Redis is commonly used?

- Fast (in-memory), supports expiry (TTL) & lightweight.

Working :- 1. Request comes to server A

2. Server checks Redis $\rightarrow$ increment counter $\rightarrow$ checks value

3. If value $>$ limit $\rightarrow$ reject request

Note:- Since Redis is shared, all servers see the same count.

Advanced concepts

Large systems may use:-

- Shared Redis clusters
- Edge level rate limiting (CDN)
- Adaptive rate limiting (based on load)
- Hierarchical rate limits.